

Time Series Analysis

- Data collected in a time sequence (daily, weekly, monthly).

Examples:

- Stock prices 📈 (daily)
- Retail sales 🛒 (monthly)
- Temperature 🌡️ (hourly)

Helps understand patterns (trend, seasonality) and predict future values.

Import & Setup

```
import pandas as pd
import matplotlib.pyplot as plt
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima.model import ARIMA
from pandas.plotting import autocorrelation_plot
```

Date Handling

Convert to datetime

```
df['Order_Date'] = pd.to_datetime(df['Order_Date'])
```

Set as index for time-based operations

```
df = df.set_index('Order_Date')
```

Extract features

```
df['Year'] = df.index.year
```

```
df['Month'] = df.index.month
```

```
df['Weekday'] = df.index.day_name()
```

```
print(df.head())
```

Makes data time-aware so you can group/slice easily.

	Product_ID	Sale_Date	Sales_Rep	Region	Sales_Amount	Quantity_Sold	Product_Category	Unit_Cost	Unit_Price	Customer_Type	Discount	Payment_Method
0	1052	2023-02-03	Bob	North	5053.97	18	Furniture	152.75	267.22	Returning	0.09	Cash
1	1093	2023-04-21	Bob	West	4384.02	17	Furniture	3816.39	4209.44	Returning	0.11	Cash
2	1015	2023-09-21	David	South	4631.23	30	Food	261.56	371.40	Returning	0.20	Bank Transfer
3	1072	2023-08-24	Bob	South	2167.94	39	Clothing	4330.03	4467.75	New	0.02	Credit Card
4	1061	2023-03-24	Charlie	East	3750.20	13	Electronics	637.37	692.71	New	0.08	Credit Card

Filtering & Slicing (Simple)

```
df['2023'] # All 2023 data
```

```
df['2023-07'] # July 2023
```

```
df['2023-07-01':'2023-07-15'] # First 15 days of July
```

Very useful in business reports.

	Product_ID	Sale_Date	Sales_Rep	Region	Sales_Amount	\
Order_Date						
2023-02-03	1052	2023-02-03	Bob	North	5053.97	
2023-04-21	1093	2023-04-21	Bob	West	4384.02	
2023-09-21	1015	2023-09-21	David	South	4631.23	
2023-08-24	1072	2023-08-24	Bob	South	2167.94	
2023-03-24	1061	2023-03-24	Charlie	East	3750.20	

	Quantity_Sold	Product_Category	Unit_Cost	Unit_Price	\
Order_Date					
2023-02-03	18	Furniture	152.75	267.22	
2023-04-21	17	Furniture	3816.39	4209.44	
2023-09-21	30	Food	261.56	371.40	
2023-08-24	39	Clothing	4330.03	4467.75	
2023-03-24	13	Electronics	637.37	692.71	

	Customer_Type	Discount	Payment_Method	Sales_Channel	\
Order_Date					
2023-02-03	Returning	0.09	Cash	Online	
2023-04-21	Returning	0.11	Cash	Retail	
2023-09-21	Returning	0.20	Bank Transfer	Retail	
2023-08-24	New	0.02	Credit Card	Retail	
2023-03-24	New	0.08	Credit Card	Online	

	Region_and_Sales_Rep	Year	Month	Weekday
Order_Date				
2023-02-03	North-Bob	2023	2	Friday
2023-04-21	West-Bob	2023	4	Friday
2023-09-21	South-David	2023	9	Thursday
2023-08-24	South-Bob	2023	8	Thursday
2023-03-24	East-Charlie	2023	3	Friday

Resampling & Aggregation

Daily sales

```
daily_sales = df['Sales'].resample('D').sum()
```

Weekly profit

```
weekly_profit = df['Profit'].resample('W').mean()
```

Monthly sales trend

```
monthly_sales = df['Sales'].resample('M').sum()
```

Converts irregular data → meaningful summaries.

```
Order_Date
2023-01-01    -159489.45666
2023-01-08    -84268.10217
2023-01-15   -153331.53100
2023-01-22   -130210.11833
2023-01-29   -96882.90708
2023-02-05   -97360.57615
2023-02-12   -117043.98178
2023-02-19   -122710.23800
2023-02-26   -140670.20538
2023-03-05   -122007.71040
2023-03-12   -110404.57500
2023-03-19   -144755.77700
2023-03-26   -91785.06210
2023-04-02   -144229.44809
2023-04-09   -106290.28500
2023-04-16   -131194.30500
2023-04-23   -155851.29529
2023-04-30   -130951.28000
2023-05-07   -183020.90300
2023-05-14   -132528.61681
2023-05-21   -187257.70230
2023-05-28   -111243.35550
2023-06-04   -99668.58375
```


Rolling Statistics (Trend Smoothing)

7-day moving average

```
df['Sales_MA7'] = df['Sales'].rolling(7).mean()
```

```
df[['Sales','Sales_MA7']].plot(figsize=(10,5), title="Sales with 7-Day Moving Avg")
plt.show()
```

Removes noise, highlights underlying trend.



Exponential Moving Average (Faster Reaction)

```
df['Sales_EMA'] = df['Sales'].ewm(span=12, adjust=False).mean()
```

Unlike simple MA, EMA reacts quicker to recent changes.

```
Order_Date
2023-02-03    5053.970000
2023-04-21    4950.900769
2023-09-21    4901.720651
2023-08-24    4481.139012
2023-03-24    4368.686857
...
2023-04-15    4497.510844
2023-09-07    4531.179945
2023-04-27    5007.875338
2023-12-20    4488.120670
2023-08-16    4555.168260
Name: Sales_EMA, Length: 1000, dtype: float64
```

Growth Analysis (MoM & YoY)

Month-over-Month %

```
df['MoM'] = df['Sales'].pct_change( periods=1 ) * 100 # Year-over-Year %
```

```
df['YoY'] = df['Sales'].pct_change( periods=12 ) * 100
```

Key business KPI for tracking growth.

```
Order_Date
2023-02-03      NaN
2023-04-21    -13.255916
2023-09-21     5.638889
2023-08-24    -53.188678
2023-03-24    72.984492
...
2023-04-15    -47.942156
2023-09-07     -0.370098
2023-04-27     61.770942
2023-12-20    -78.643066
2023-08-16    202.179850
Name: MoM, Length: 1000, dtype: float64
```

Seasonal Decomposition (Trend + Seasonality)


```
result = seasonal_decompose(df['Sales'], model='additive', period=12)
```

```
result.plot()
```

```
plt.show()
```

Breaks series into:

- Trend → long-term growth
- Seasonality → repeating cycles (holidays, seasons)
- Residual → noise



Cumulative Sum

```
df['Cumulative_Sales'] = df['Sales'].cumsum()
```

Tracks total growth over time.

```
Order_Date
2023-04-15    5000365.77
2023-09-07    5005082.13
2023-04-27    5012711.83
2023-12-20    5014341.30
2023-08-16    5019265.23
Name: Cumulative_Sales, dtype: float64
```

Lag Features (Feature Engineering for ML)

```
df['Sales_Lag1'] = df['Sales'].shift(1) # yesterday's sales
```

```
df['Sales_Lag7'] = df['Sales'].shift(7) # last week's sales
```


Helps models learn from past values.

```
Order_Date
2023-02-03      NaN
2023-04-21      NaN
2023-09-21      NaN
2023-08-24      NaN
2023-03-24      NaN
...
2023-04-15    4912.69
2023-09-07    9215.32
2023-04-27     496.59
2023-12-20    2985.46
2023-08-16    2154.66
Name: Sales_Lag7, Length: 1000, dtype: float64
```

Rolling Correlation (Sales vs Profit)

```
df['Sales'].rolling(30).corr(df['Profit']).plot(title="Rolling Correlation (Sales vs Profit)")  
plt.show()
```

Measures relationship strength over time.



Stationarity Test (Needed for Forecasting)

```
from statsmodels.tsa.stattools import adfuller
```

```
result = adfuller(df['Sales'])  
print("ADF Statistic:", result[0])  
print("p-value:", result[1])
```


If $p < 0.05 \rightarrow$ stationary (good for ARIMA).

ADF Statistic: -30.416398670618502
p-value: 0.0

Autocorrelation (Check Repetition)

```
autocorrelation_plot(df['Sales'])  
plt.show()
```


Shows seasonal lags (e.g., sales spike every 7 days).



Forecasting with ARIMA (Basic)

```
model = ARIMA(df['Sales'], order=(1,1,1))  
fit = model.fit()  
  
df['Forecast'] = fit.predict(start=len(df)-12, end=len(df)+6, dynamic=True)  
  
df[['Sales', 'Forecast']].plot(figsize=(10,5), title="Sales Forecast with ARIMA")  
plt.show()
```

Predicts future values based on past trends.



Train-Test Forecast Validation

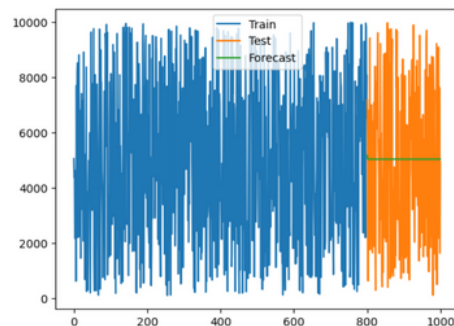
```
train = df['Sales'][:int(0.8*len(df))]  
test = df['Sales'][int(0.8*len(df)):]
```

```
model = ARIMA(train, order=(1,1,1))  
fit = model.fit()
```

```
forecast = fit.predict(start=len(train), end=len(df)-1, dynamic=False)
```

```
plt.plot(train, label="Train")  
plt.plot(test, label="Test")  
plt.plot(forecast, label="Forecast")  
plt.legend()  
plt.show()
```


Compares predicted vs actual.



Seasonal Plot (Month vs Avg Sales)

```
df.groupby(df.index.month)['Sales'].mean().plot(kind='bar')  
plt.title("Average Sales by Month")  
plt.show()
```

Identifies best & worst sales months.



Heatmap of Seasonality

```
import seaborn as sns
```

```
pivot = df.pivot_table(values="Sales", index=df.index.year, columns=df.index.month, aggfunc='sum')
sns.heatmap(pivot, annot=True, fmt=".Of", cmap="YlGnBu")
plt.title("Year-Month Sales Heatmap")
plt.show()
```

Visualizes seasonal patterns across years.

